

Building a Bayesian Burrito:

John McDonnell, MD

Description

San Diego has long been widely celebrated for its prime location, breezy ocean views, and fine dining experiences. One thing that only the locals fully appreciate is how great its Mexican food offerings are— no surprise given that the Mexican border is only 15 miles away.

As a burrito aficionado (I gained a remarkable 15 pounds during my first year in medical school due to a diet of [almost] exclusively Chipotle burritos) and former denizen of San Diego, when I saw that a like-minded searcher from the San Diego area had collected years of data on the subject of local burrito options in San Diego I grew excited— and a little bit hungry...

Data Source

This project uses data from [Scott Cole's Github site](#). A few years ago, Scott and some San Diego friends started rating local burritos on a 10 point scale. In his words:

1. "Volume - "size matters," "bigger is better," or whatever your favorite innuendo is fits because there's nothing more disconcerting than ordering a burrito and not being full.
2. Tortilla quality
3. Temperature - the Goldilocks zone
4. Meat quality
5. Non-meat filling quality
6. Meat : filling - The ratio between meat and non-meat. Perhaps the golden ratio: 1.6180339887...
7. Uniformity - Bites full of sour cream and cheese with no meat are disappointing.
8. Salsa quality - and variety!
9. Flavor synergy - "That magical aspect a great burrito has, making everything come together like it is a gift from the skies" - A wise Dutchman
10. Wrap integrity - you ordered a burrito, not a burrito bowl."

Based on the above 10 point scale, Scott and his burrito-loving friends also rendered a final burrito rating (on a 5 point, noninteger scale).

Study Population

Our study population is the burritos of the greater San Diego area. Mr. Cole and colleagues rated 423 burritos from May 2011 to August 2019. In addition to the 10 point ratings we saw before, he also collected data on burrito ingredients.

Variables

A description of the data and the variables of interest is seen below.

```
. describe
```

Contains data from burrito_clean.dta

```
obs:      423
```

```
vars:      60
```

```
23 Jul 2020 09:27
```

variable name	storage type	display format	value label	variable label
burrito	str30	%30s		burrito name
date	str10	%10s		date of rating
neighborhood	str18	%18s		neighborhood in San Diego
yelp	float	%9.0g		yelp rating
google	float	%9.0g		google rating
cost	float	%9.0g		cost of burrito (\$)
hunger	float	%9.0g		hunger level of reviewer
massg	int	%8.0g		mass of burrito (g)
densitygml	float	%9.0g		density of burrito (g/mL)
length	float	%9.0g		length of burrito
circum	float	%9.0g		circumference of burrito
volume	float	%9.0g		volume of burrito
tortilla	float	%9.0g		rating of quality of tortilla
temp	float	%9.0g		rating of temperature of
burrito				
meat	float	%9.0g		rating of meat in burrito
fillings	float	%9.0g		rating of non-meat fillings in
the burrito				
meatfilling	float	%9.0g		ratio between meat and non-meat
uniformity	float	%9.0g		rating of uniformity of burrito
salsa	float	%9.0g		rating of salsa used in burrito
synergy	float	%9.0g		rating of the burrito's synergy
wrap	float	%9.0g		integrity of burrito's wrap job
overall	float	%9.0g		overall rating of the burrito
reviewer	str12	%12s		reviewer name
date_num	float	%d		date of rating
l_location	str51	%51s		restaurant name
chips_num	float	%9.0g		indicator for chips
recommend	float	%9.0g		would reviewer recommend
burrito				
beef_ind	float	%9.0g		indicator for beef
pico_ind	float	%9.0g		indicator for pico
guac_ind	float	%9.0g		indicator for guac
cheese_ind	float	%9.0g		indicator for cheese
fries_ind	float	%9.0g		indicator for fries
sourcream_ind	float	%9.0g		indicator for sourcream
pork_ind	float	%9.0g		indicator for pork
chicken_ind	float	%9.0g		indicator for chicken
shrimp_ind	float	%9.0g		indicator for shrimp
fish_ind	float	%9.0g		indicator for fish
rice_ind	float	%9.0g		indicator for rice

beans_ind	float	%9.0g	indicator for beans
lettuce_ind	float	%9.0g	indicator for lettuce
tomato_ind	float	%9.0g	indicator for tomato
bellpeper_ind	float	%9.0g	indicator for bellpeper
carrots_ind	float	%9.0g	indicator for carrots
cabbage_ind	float	%9.0g	indicator for cabbage
sauce_ind	float	%9.0g	indicator for sauce
cilantro_ind	float	%9.0g	indicator for cilantro
onion_ind	float	%9.0g	indicator for onion
taquito_ind	float	%9.0g	indicator for taquito
pineapple_ind	float	%9.0g	indicator for pineapple
ham_ind	float	%9.0g	indicator for ham
chilerelleno_ind	float	%9.0g	indicator for chilerelleno
nopales_ind	float	%9.0g	indicator for nopales
lobster_ind	float	%9.0g	indicator for lobster
egg_ind	float	%9.0g	indicator for egg
mushroom_ind	float	%9.0g	indicator for mushroom
bacon_ind	float	%9.0g	indicator for bacon
sushi_ind	float	%9.0g	indicator for sushi
avocado_ind	float	%9.0g	indicator for avocado
corn_ind	float	%9.0g	indicator for corn
zucchini_ind	float	%9.0g	indicator for zucchini

Sorted by:

Note: Dataset has changed since last saved.

I've attached labels to each variable to specify its meaning in a clear way.

We will go into specifics in a minute, but for now let's examine the string/character variables. There are five of them in total: *burrito*, *date*, *neighborhood*, *reviewer*, and *l_location*. These variables are not going to be very useful to us in the regression analysis, but we can get a "flavor" of them (pun most certainly intended) by looking at a small subset.

```
. list burrito date neighborhood reviewer l_location in 1/5
```

```

+-----+
--|
|      burrito      date  neighbor~d  reviewer      l_location
|
|-----|
-|
1. |  california  1/18/2016      miramar      Scott      donato's taco shop
|
2. |  california  1/24/2016    san marcos      Scott      oscar's mexican food
|
3. |    carnitas  1/24/2016                Emily      oscar's mexican food
|
4. | carne asada  1/24/2016                Ricardo     oscar's mexican food
|
5. |  california  1/27/2016      carlsbad      Scott      pollos maria
|

```

```

+-----+
-+

```

For our analysis, we are definitely going to need to look at the continuous variables. Many of these correspond to the reviewers' 1-5 rating scale for individual elements of burrito quality, while others pertain to burrito measurements. As we will see later, there are a lot of missing data with the objective burrito measurements but most of the subjective ratings are complete. Additionally, we see here variables corresponding to the burritos' Yelp and Google reviews; however, many of these are missing.

```
. summarize yelp-overall
```

Variable	Obs	Mean	Std. Dev.	Min	Max
yelp	87	3.887356	.4753957	2.5	4.5
google	87	4.167816	.3736975	2.9	5
cost	416	7.065216	1.503645	2.99	25
hunger	420	3.496095	.8114664	.5	5
massg	22	546.1818	144.4456	350	925
densitygml	22	.6752774	.0804682	.56	.8656716
length	284	20.0469	2.084957	15	26
circum	282	22.13174	1.777526	17	29
volume	282	.7864894	.1522597	.4	1.54
tortilla	423	3.519385	.7933014	1	5
temp	403	3.780397	.9800436	1	5
meat	409	3.622249	.8283836	1	5
fillings	420	3.542024	.8012532	1	5
meatfilling	414	3.589082	.9962922	.5	5
uniformity	421	3.434086	1.069349	0	5
salsa	398	3.372613	.9224345	0	5
synergy	421	3.587767	.886277	1	5
wrap	420	3.98119	1.115803	0	5
overall	421	3.620887	.755718	1	5

Finally, we also have a number of indicator variables, generally corresponding to whether a given ingredient was included in the burrito. As we might expect, when we look at the summarized proportions (I included them this way to spare the reader a long list of tabulated output), we can see that certain ingredients were common (guacamole, beef) and others were uncommon (lobster).

```
. codebook chips_num-zucchini_ind, compact
```

Variable	Obs	Unique	Mean	Min	Max	Label
chips_num	423	2	.0520095	0	1	indicator for chips
recommend	233	2	.6995708	0	1	would reviewer recommend
burrito						

beef_ind	423	2	.4255319	0	1	indicator for beef
pico_ind	423	2	.3758865	0	1	indicator for pico
guac_ind	423	2	.3664303	0	1	indicator for guac
cheese_ind	423	2	.3782506	0	1	indicator for cheese
fries_ind	423	2	.3026005	0	1	indicator for fries
sourcream_~d	423	2	.2174941	0	1	indicator for sourcream
pork_ind	423	2	.1205674	0	1	indicator for pork
chicken_ind	423	2	.0496454	0	1	indicator for chicken
shrimp_ind	423	2	.0496454	0	1	indicator for shrimp
fish_ind	423	2	.0141844	0	1	indicator for fish
rice_ind	423	2	.0851064	0	1	indicator for rice
beans_ind	423	2	.0827423	0	1	indicator for beans
lettuce_ind	423	2	.0260047	0	1	indicator for lettuce
tomato_ind	423	2	.0165485	0	1	indicator for tomato
bellpeper_~d	423	2	.0165485	0	1	indicator for bellpeper
carrots_ind	423	2	.0023641	0	1	indicator for carrots
cabbage_ind	423	2	.0189125	0	1	indicator for cabbage
sauce_ind	423	2	.0898345	0	1	indicator for sauce
cilantro_ind	423	2	.035461	0	1	indicator for cilantro
onion_ind	423	2	.0401891	0	1	indicator for onion
taquito_ind	423	2	.0094563	0	1	indicator for taquito
pineapple_~d	423	2	.0165485	0	1	indicator for pineapple
ham_ind	423	2	.0047281	0	1	indicator for ham
chilerelle~d	423	2	.0094563	0	1	indicator for chilerelleno
nopales_ind	423	2	.0094563	0	1	indicator for nopales
lobster_ind	423	2	.0023641	0	1	indicator for lobster
egg_ind	423	2	.0118203	0	1	indicator for egg
mushroom_ind	423	2	.0070922	0	1	indicator for mushroom
bacon_ind	423	2	.0070922	0	1	indicator for bacon
sushi_ind	423	2	.0047281	0	1	indicator for sushi
avocado_ind	423	2	.0307329	0	1	indicator for avocado
corn_ind	423	2	.0070922	0	1	indicator for corn
zucchini_ind	423	2	.0023641	0	1	indicator for zucchini

Response Variable

Our response variable in this analysis is *top_burrito*. This variable is a binary corresponding to 1 if a burrito was ranked in the top 2 categories for overall rating (*overall_ord*).

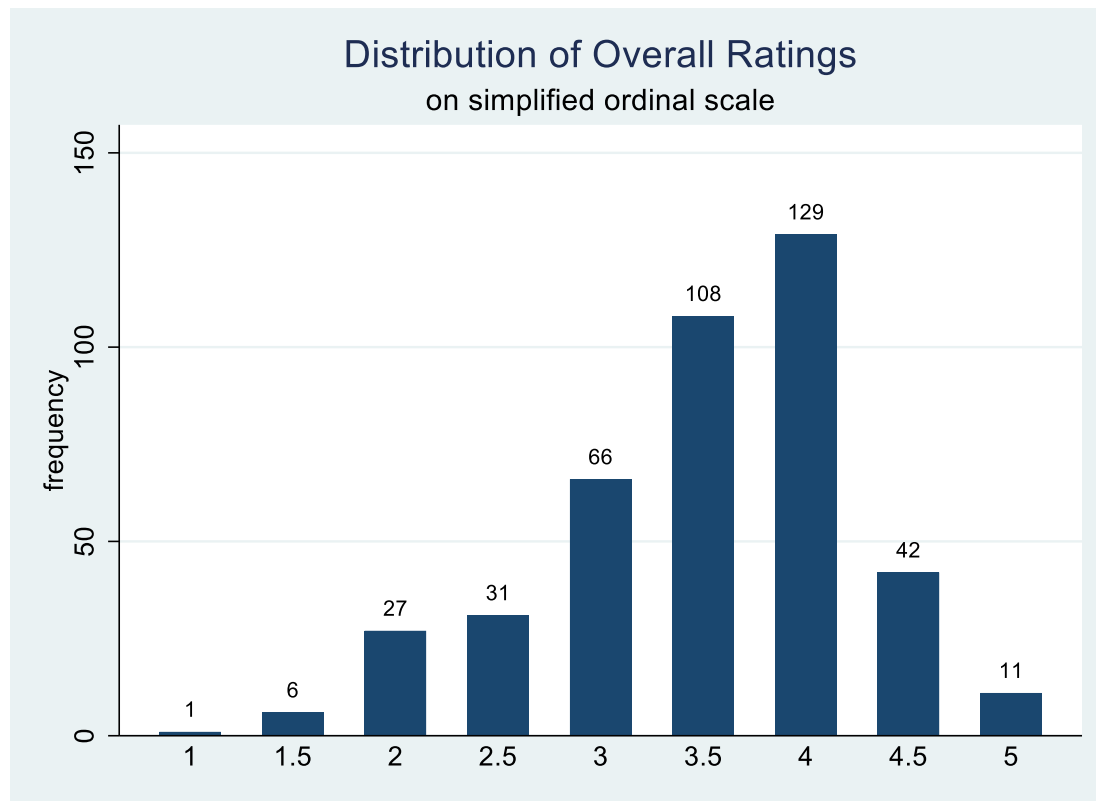
The distribution of the “overall” burrito rating is given below, first in table form (note that “.” indicates missing in Stata):

```
. tab overall_ord, missing
```

overall_ord	Freq.	Percent	Cum.
1	1	0.24	0.24
1.5	6	1.42	1.65
2	27	6.38	8.04
2.5	31	7.33	15.37
3	66	15.60	30.97

3.5		108	25.53	56.50
4		129	30.50	87.00
4.5		42	9.93	96.93
5		11	2.60	99.53
.		2	0.47	100.00
<hr/>				
Total		423	100.00	

and then in graphical form (focusing on the bars corresponding to ratings of 4, 4.5, and 5):



To be honest, I would have vastly preferred to treat overall rating as an ordinal outcome variable, but ran into some catastrophic issues with convergence when I fit the Bayesian ordinal regression model. These problems disappeared when I treated the outcome as a binary and ran Bayesian logistic regression instead (of course, at the expense of some loss of information in the outcome). In any case, our outcome variable, *top_burrito*, will simply be binary for all the variables that were rated as 4 or higher.

```
. gen top_burrito = 0

. replace top_burrito = 1 if overall_ord >=4 & !missing(overall_ord)
(182 real changes made)

.
. tab overall_ord top_burrito

overall_or |      top_burrito
```

d	0	1	Total
1	1	0	1
1.5	6	0	6
2	27	0	27
2.5	31	0	31
3	66	0	66
3.5	108	0	108
4	0	129	129
4.5	0	42	42
5	0	11	11
Total	239	182	421

Predictors

While a number of predictor variables will be considered here, we are ultimately going to have to choose just a few of them in the final model. We will spend a lot more time discussing our variable selection process a little bit later, but here is a “sneak peak” of the predictor variables we will end up using.

```
. list synergy meat meatfilling fillings in 1/5
```

	synergy	meat	meatfi~g	fillings
1.	4	3	4	3.5
2.	2.5	2.5	2	2.5
3.	3	2.5	4.5	3
4.	4	3.5	4	3
5.	4.5	4	4.5	3.5

Objective of Study

What makes a burrito good? I will use Bayesian logistic regression methods to answer this question. In this dataset, with $N = 421$ observations and a pool of over 60 candidate predictor variables, I will follow the overall strategy:

1. determine reasonable predictor variables using stepwise logistic regression and lasso methods
2. fit a Bayesian logistic regression model
3. examine fit characteristics of the Bayesian model and diagnose any problems with autocorrelation and convergence
4. make some predictions
5. compare to a standard/frequentist logistic regression model and summarize results

Background: What is known on this subject?

Perhaps unsurprisingly, there is a dearth of good information on what makes a good burrito. The most interesting previous research on the subject comes from [Five Thirty Eight](#). Nate Silver and his co-investigator Anna Barry-Jester looked for the best burrito in America with their project [In Search of America's Best Burrito](#). Their process involved data mining Yelp to create an overall rating called “Value Over Replacement Burrito.” forming a “Burrito Selection Committee” to determine the most promising 64 burrito joints (from 60k+ ratings) in the United States, and then taste-testing each of these. The ultimate winner was La Taqueria’s (San Francisco, CA) carnitas burrito whose “bombardment of liquid and flavor ... are enough to stop any woman in her tracks, even one who’d been eating burritos daily for two months straight.” Although I suspect I may have an appetite sufficient to replicate this task, I don’t think it will fit into my school, work, and home life at this point in time.

In his own analysis of the burrito data, Scott Cole performed a principal components analysis predicting rating as an outcome variable, and found fillings, meat, meat : filling, and uniformity to be the most important predictors. The details of his analysis are summarized [here](#).

Contribution to the literature

To my knowledge, this will be the first study evaluating predictors of burrito quality using Bayesian methods.

Exploratory Analysis

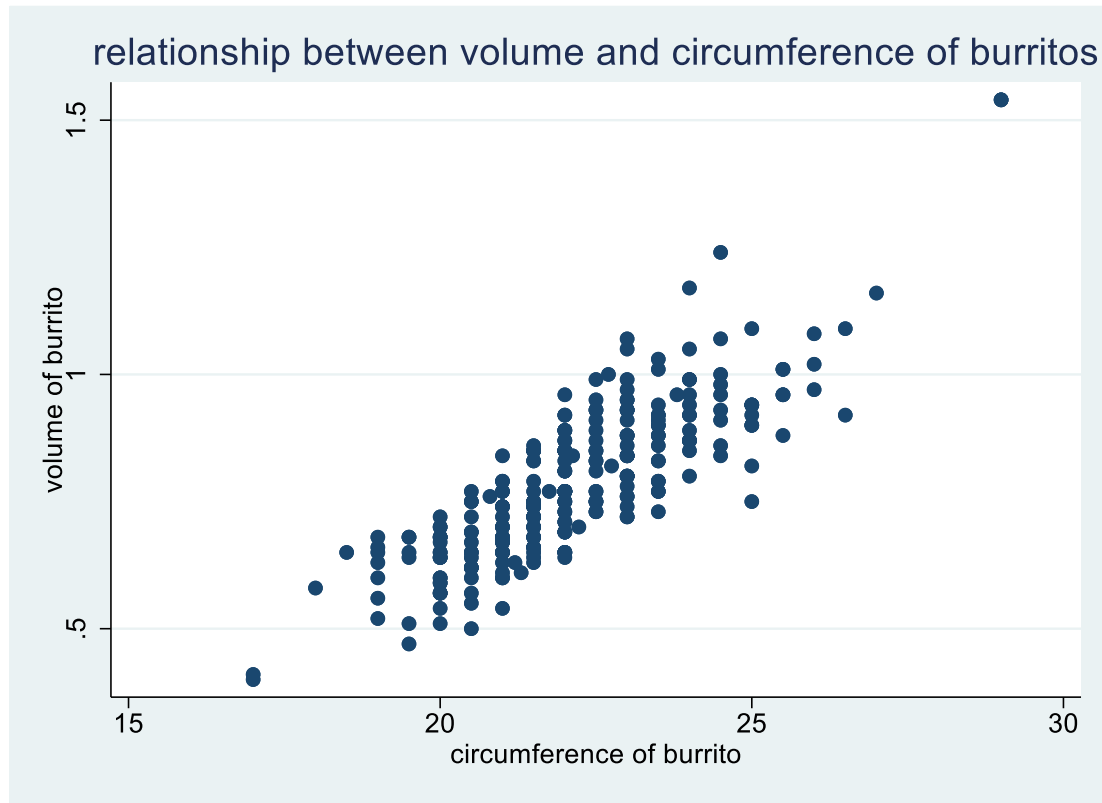
Having spoken a bit about the background and objectives to this analysis, let’s look more into the dataset itself, and particularly how our outcome of interest relates to the potential predictor variables.

A first step is to examine some scatterplots of the relationship between continuous variables. This approach can give us an overall sense of the distribution and correlations of our variables, and perhaps clue us into relationships that may be important for our model.

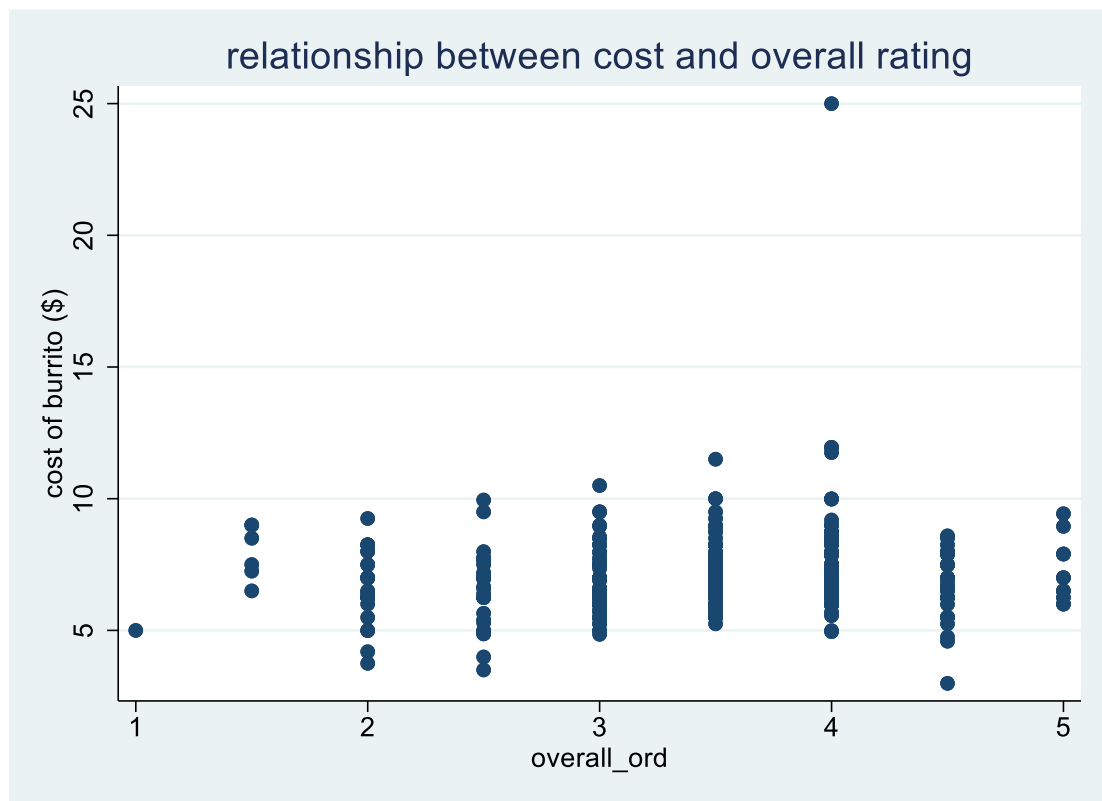
For the purposes of my exploratory data analysis, I ran the following command (output omitted to spare the reader’s sanity):

```
graph matrix yelp google cost hunger massg densitygml length circum volume  
tortilla temp meat fillings meatfilling uniformity salsa synergy wrap  
overall_ord
```

After running this command, the first thing I noted was that the strongest linear relationships in these data were just where I would expect them– in the measures of the physical dimensions of the burritos. For example:




The relationship between the overall rating and the cost seems less strong. See:




(Incidentally, I tracked down the outlier in this plot. It turns out that this [was not a miscode](#) - scroll to the “specials” section for more on this.)

There turned out to be a monotonic relationship between the burrito’s synergy and its overall rating. This makes intuitive sense - a reviewer who thought a burrito had great synergy would be likely to rate it higher.



Another noteworthy relationship is seen here, this time between the fillings and the overall rating of the burrito.



Unfortunately, we have a problem that I alluded to before— namely, many of the continuous variables we might be interested in have a lot of missing data. One reasonable strategy would be to impute, but for the purposes of focusing on the main goals of this project, I will just leave these variables out instead.

```
. inspect massg-volume
```

massg: mass of burrito (g)				Number of Observations		
				Total	Integers	
Nonintegers						
	#		Negative	-	-	
	#	#	Zero	-	-	
	#	#	Positive	22	22	
	#	#				
	#	#	Total	22	22	
	#	#	Missing	401		
+						
350		925		423		
(18 unique values)						

densitygml: density of burrito (g/mL)	Number of Observations
---------------------------------------	------------------------

```

-----
-
Total      Integers
Nonintegers
| #      #      Negative      -      -
-
| #      #      #      Zero      -      -
-
| #      #      #      #      Positive      22      -
22
| #      #      #      #
-
| #      #      #      #      Total      22      -
22
| #      #      #      #      #      Missing      401
+-----
.56      .8656716      423
(21 unique values)

```

```

length:  length of burrito
-----
Number of Observations
-----
Total      Integers
Nonintegers
|      #      Negative      -      -
-
|      #      Zero      -      -
-
|      #      #      Positive      284      162
122
|      #      #      #
-
|      #      #      #      Total      284      162
122
| #      #      #      #      .      Missing      139
+-----
15      26      423
(29 unique values)

```

```

circum:  circumference of burrito
-----
Number of Observations
-----
Total      Integers
Nonintegers
|      #      Negative      -      -
-
|      #      #      Zero      -      -
-
|      #      #      Positive      282      162
120
|      #      #
-
|      #      #      Total      282      162
120
| .      #      #      #      .      Missing      141
+-----
17      29      423

```

(30 unique values)

volume: volume of burrito				Number of Observations		
				Total	Integers	
Nonintegers						
	#		Negative	-	-	
	#		Zero	-	-	
	#		Positive	282	2	
280	#					
	#					
	#	#	Total	282	2	
280	#	#				
	#	#	Missing	141		
+	#	#				
.4		1.54		423		
(64 unique values)						

. drop massg-volume

The variables corresponding to Yelp and Google reviews also have a lot of missing values. Additionally, these might not be good predictors for another reason– they would be expected to be highly correlated with the outcome rating. Therefore, we will leave them out.

. inspect yelp-google

yelp: yelp rating				Number of Observations		
				Total	Integers	
Nonintegers						
	#		Negative	-	-	
	#		Zero	-	-	
	#		Positive	87	48	
39	#					
	#					
	#	#	Total	87	48	
39	#	#				
	.	#	Missing	336		
+	.	#				
2.5		4.5		423		
(6 unique values)						

google: google rating				Number of Observations		
				Total	Integers	
Nonintegers						

```

|          #          Negative          -          -
-
|          #  #          Zero          -          -
-
|          #  #          Positive          87          10
77
|          #  #          -----
-
|          #  #          Total          87          10
77
| . . # # #          Missing          336
+-----+-----+
2.9          5          423
(18 unique values)

. drop yelp-google

```

Let's explore the categorical variables now. Perhaps the best way to get a sense of the relationships at play here is by using tables. For example, we can see that many of the "top burritos" contained beef, but not the majority.

```

. tab top_burrito beef, row

+-----+
| Key          |
|-----|
| frequency    |
| row percentage |
+-----+

top_burrit | indicator for beef
o | 0 1 | Total
-----+-----+-----+
0 | 128 113 | 241
| 53.11 46.89 | 100.00
-----+-----+-----+
1 | 115 67 | 182
| 63.19 36.81 | 100.00
-----+-----+-----+
Total | 243 180 | 423
| 57.45 42.55 | 100.00

```

In fact, a large number of these fillings were rare, with counts ≤ 5 . For this reason, we won't seriously consider them as predictors. I'll also drop any indicator variables where the "d" cell (in standard epidemiological format, corresponding in this case to counts where $\text{top_burrito} == 1$ and $\text{topping} == 1$) is ≤ 5 . The other indicators, however, are fair game for possible inclusion.

```

.
. drop zucchini corn sushi bacon mushroom egg lobster nopales chile ham
taquito carrot /* low counts overall */

. drop avocado pineapple onion cilantro cabbage bellpeper tomato lettuce
beans fish shrimp chicken /* low counts in "d" */

```

```
.
. foreach v of varlist beef-sauce{
2.     tab top_burrito `v', row
3. }
```

```
+-----+
| Key      |
+-----+
| frequency |
| row percentage |
+-----+
```

top_burrit		indicator for beef		
	o	0	1	Total
0		128	113	241
		53.11	46.89	100.00
1		115	67	182
		63.19	36.81	100.00
Total		243	180	423
		57.45	42.55	100.00

```
+-----+
| Key      |
+-----+
| frequency |
| row percentage |
+-----+
```

top_burrit		indicator for pico		
	o	0	1	Total
0		137	104	241
		56.85	43.15	100.00
1		127	55	182
		69.78	30.22	100.00
Total		264	159	423
		62.41	37.59	100.00

```
+-----+
| Key      |
+-----+
| frequency |
| row percentage |
+-----+
```

top_burrit		indicator for guac		
	o	0	1	Total
0		149	92	241
		61.83	38.17	100.00

1	119	63	182
	65.38	34.62	100.00
<hr/>			
Total	268	155	423
	63.36	36.64	100.00

```

+-----+
| Key    |
+-----+
| frequency |
| row percentage |
+-----+

```

top_burrit indicator for cheese			
o	0	1	Total
<hr/>			
0	145	96	241
	60.17	39.83	100.00
<hr/>			
1	118	64	182
	64.84	35.16	100.00
<hr/>			
Total	263	160	423
	62.17	37.83	100.00

```

+-----+
| Key    |
+-----+
| frequency |
| row percentage |
+-----+

```

top_burrit indicator for fries			
o	0	1	Total
<hr/>			
0	166	75	241
	68.88	31.12	100.00
<hr/>			
1	129	53	182
	70.88	29.12	100.00
<hr/>			
Total	295	128	423
	69.74	30.26	100.00

```

+-----+
| Key    |
+-----+
| frequency |
| row percentage |
+-----+

```

top_burrit indicator for sourcream			
o	0	1	Total
<hr/>			
0	188	53	241
	78.01	21.99	100.00

1	143	39	182
	78.57	21.43	100.00
Total	331	92	423
	78.25	21.75	100.00

```

+-----+
| Key    |
+-----+
| frequency |
| row percentage |
+-----+

```

top_burrit indicator for pork			
o	0	1	Total
0	208	33	241
	86.31	13.69	100.00
1	164	18	182
	90.11	9.89	100.00
Total	372	51	423
	87.94	12.06	100.00

```

+-----+
| Key    |
+-----+
| frequency |
| row percentage |
+-----+

```

top_burrit indicator for rice			
o	0	1	Total
0	213	28	241
	88.38	11.62	100.00
1	174	8	182
	95.60	4.40	100.00
Total	387	36	423
	91.49	8.51	100.00

```

+-----+
| Key    |
+-----+
| frequency |
| row percentage |
+-----+

```

top_burrit indicator for sauce			
o	0	1	Total
0	216	25	241
	89.63	10.37	100.00

1	169	13	182
	92.86	7.14	100.00
Total	385	38	423
	91.02	8.98	100.00

For our remaining indicator variables, there are none with an obviously strong relationship to the outcome.

Model Selection

After excluding the variables mentioned above, we are left with the following:

```
. describe burrito-sauce
```

variable name	storage type	display format	value label	variable label

burrito	str30	%30s		burrito name
date	str10	%10s		date of rating
neighborhood	str18	%18s		neighborhood in San Diego
cost	float	%9.0g		cost of burrito (\$)
hunger	float	%9.0g		hunger level of reviewer
tortilla	float	%9.0g		rating of quality of tortilla
temp	float	%9.0g		rating of temperature of
burrito				
meat	float	%9.0g		rating of meat in burrito
fillings	float	%9.0g		rating of non-meat fillings in
the burrito				
meatfilling	float	%9.0g		ratio between meat and non-meat
uniformity	float	%9.0g		rating of uniformity of burrito
salsa	float	%9.0g		rating of salsa used in burrito
synergy	float	%9.0g		rating of the burrito's synergy
wrap	float	%9.0g		integrity of burrito's wrap job
overall	float	%9.0g		overall rating of the burrito
reviewer	str12	%12s		reviewer name
date_num	float	%d		date of rating
l_location	str51	%51s		restaurant name
chips_num	float	%9.0g		indicator for chips
recommend	float	%9.0g		would reviewer recommend
burrito				
beef_ind	float	%9.0g		indicator for beef
pico_ind	float	%9.0g		indicator for pico
guac_ind	float	%9.0g		indicator for guac
cheese_ind	float	%9.0g		indicator for cheese
fries_ind	float	%9.0g		indicator for fries
sourcream_ind	float	%9.0g		indicator for sourcream
pork_ind	float	%9.0g		indicator for pork
rice_ind	float	%9.0g		indicator for rice
sauce_ind	float	%9.0g		indicator for sauce

I will employ two strategies for model selection. The first is stepwise forward regression. Note that Stata accomplishes this using changes in p values, not an information criterion like AIC or BIC.

```
. set seed 453
```

```
.
. stepwise, pe(.2): logit top_burrito c.cost c.hunger c.(tortilla-wrap)
i.(chips beef-sauce)
```

```
note: 0b.chips_num dropped because of estimability
note: 0b.beef_ind dropped because of estimability
note: 0b.pico_ind dropped because of estimability
note: 0b.guac_ind dropped because of estimability
note: 0b.cheese_ind dropped because of estimability
note: 0b.fries_ind dropped because of estimability
note: 0b.sourcream_ind dropped because of estimability
note: 0b.pork_ind dropped because of estimability
note: 0b.rice_ind dropped because of estimability
note: 0b.sauce_ind dropped because of estimability
```

```
begin with empty model
```

```
p = 0.0000 < 0.2000 adding synergy
p = 0.0000 < 0.2000 adding meat
p = 0.0000 < 0.2000 adding meatfilling
p = 0.0002 < 0.2000 adding temp
p = 0.0007 < 0.2000 adding fillings
p = 0.0035 < 0.2000 adding tortilla
p = 0.0532 < 0.2000 adding salsa
p = 0.1975 < 0.2000 adding 1.rice_ind
p = 0.1991 < 0.2000 adding cost
```

Logistic regression	Number of obs	=
358		
	LR chi2(9)	=
278.23		
	Prob > chi2	=
0.0000		
Log likelihood = -103.6348	Pseudo R2	=
0.5731		

top_burrito	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
synergy	1.716266	.4052795	4.23	0.000	.9219331
meat	1.699645	.3629123	4.68	0.000	.9883498
meatfilling	1.069974	.2472165	4.33	0.000	.5854387
temp	.7140077	.1992907	3.58	0.000	.3234051
fillings	1.039525	.3543918	2.93	0.003	.3449296

tortilla	.8013605	.279028	2.87	0.004	.2544756	
1.348245						
salsa	.3800812	.2372236	1.60	0.109	-.0848686	
.8450309						
1.rice_ind	-.9454823	.6962917	-1.36	0.175	-2.310189	
.4192243						
cost	.2000829	.1558271	1.28	0.199	-.1053326	
.5054985						
_cons	-29.16476	3.500084	-8.33	0.000	-36.0248	-
22.30473						

-

.

Our stepwise selection technique suggests we include *synergy*, *meat*, *meatfilling*, *temp*, *fillings*, *tortilla*, *salsa*, *cost* and *rice_ind* in the model.

The other strategy for variable selection I want to try is Lasso penalized regression.

```
. set seed 453

. lasso logit top_burrito c.cost c.hunger c.(tortilla-wrap) i.(chips beef-
sauce)

10-fold cross-validation with 100 lambdas ...
Grid value 1:      lambda = .2968057    no. of nonzero coef. =      0
Folds: 1...5....10  CVF = 1.362455
Grid value 2:      lambda = .2704383    no. of nonzero coef. =      1
Folds: 1...5....10  CVF = 1.306353
Grid value 3:      lambda = .2464133    no. of nonzero coef. =      1
Folds: 1...5....10  CVF = 1.252595
Grid value 4:      lambda = .2245226    no. of nonzero coef. =      2
Folds: 1...5....10  CVF = 1.202766
Grid value 5:      lambda = .2045767    no. of nonzero coef. =      2
Folds: 1...5....10  CVF = 1.155658
Grid value 6:      lambda = .1864026    no. of nonzero coef. =      2
Folds: 1...5....10  CVF = 1.113017
Grid value 7:      lambda = .1698431    no. of nonzero coef. =      3
Folds: 1...5....10  CVF = 1.07567
Grid value 8:      lambda = .1547548    no. of nonzero coef. =      3
Folds: 1...5....10  CVF = 1.042219
Grid value 9:      lambda = .1410068    no. of nonzero coef. =      3
Folds: 1...5....10  CVF = 1.011164
Grid value 10:     lambda = .1284801    no. of nonzero coef. =      3
Folds: 1...5....10  CVF = .9836577
Grid value 11:     lambda = .1170663    no. of nonzero coef. =      4
Folds: 1...5....10  CVF = .9579422
Grid value 12:     lambda = .1066664    no. of nonzero coef. =      4
Folds: 1...5....10  CVF = .9330114
Grid value 13:     lambda = .0971905    no. of nonzero coef. =      5
Folds: 1...5....10  CVF = .9072486
Grid value 14:     lambda = .0885564    no. of nonzero coef. =      6
Folds: 1...5....10  CVF = .8825074
Grid value 15:     lambda = .0806893    no. of nonzero coef. =      6
Folds: 1...5....10  CVF = .8591103
```


Grid value 16:	lambda = .073521	no. of nonzero coef. =	6
Folds: 1...5....10	CVF = .8381328		
Grid value 17:	lambda = .0669896	no. of nonzero coef. =	6
Folds: 1...5....10	CVF = .8178472		
Grid value 18:	lambda = .0610385	no. of nonzero coef. =	6
Folds: 1...5....10	CVF = .7995145		
Grid value 19:	lambda = .055616	no. of nonzero coef. =	7
Folds: 1...5....10	CVF = .7829686		
Grid value 20:	lambda = .0506752	no. of nonzero coef. =	7
Folds: 1...5....10	CVF = .7679245		
Grid value 21:	lambda = .0461734	no. of nonzero coef. =	7
Folds: 1...5....10	CVF = .7541		
Grid value 22:	lambda = .0420714	no. of nonzero coef. =	7
Folds: 1...5....10	CVF = .7414924		
Grid value 23:	lambda = .0383339	no. of nonzero coef. =	7
Folds: 1...5....10	CVF = .7301799		
Grid value 24:	lambda = .0349285	no. of nonzero coef. =	7
Folds: 1...5....10	CVF = .720204		
Grid value 25:	lambda = .0318255	no. of nonzero coef. =	7
Folds: 1...5....10	CVF = .7115774		
Grid value 26:	lambda = .0289982	no. of nonzero coef. =	8
Folds: 1...5....10	CVF = .7040825		
Grid value 27:	lambda = .0264221	no. of nonzero coef. =	8
Folds: 1...5....10	CVF = .697619		
Grid value 28:	lambda = .0240748	no. of nonzero coef. =	9
Folds: 1...5....10	CVF = .6920241		
Grid value 29:	lambda = .0219361	no. of nonzero coef. =	9
Folds: 1...5....10	CVF = .6874287		
Grid value 30:	lambda = .0199873	no. of nonzero coef. =	10
Folds: 1...5....10	CVF = .6844194		
Grid value 31:	lambda = .0182117	no. of nonzero coef. =	12
Folds: 1...5....10	CVF = .6818702		
Grid value 32:	lambda = .0165938	no. of nonzero coef. =	12
Folds: 1...5....10	CVF = .6799953		
Grid value 33:	lambda = .0151197	no. of nonzero coef. =	12
Folds: 1...5....10	CVF = .6787508		
Grid value 34:	lambda = .0137765	no. of nonzero coef. =	13
Folds: 1...5....10	CVF = .6781779		
Grid value 35:	lambda = .0125526	no. of nonzero coef. =	13
Folds: 1...5....10	CVF = .6776702		
Grid value 36:	lambda = .0114375	no. of nonzero coef. =	13
Folds: 1...5....10	CVF = .6772458		
Grid value 37:	lambda = .0104214	no. of nonzero coef. =	14
Folds: 1...5....10	CVF = .6771735		
Grid value 38:	lambda = .0094956	no. of nonzero coef. =	14
Folds: 1...5....10	CVF = .6774453		
Grid value 39:	lambda = .008652	no. of nonzero coef. =	14
Folds: 1...5....10	CVF = .6779057		
Grid value 40:	lambda = .0078834	no. of nonzero coef. =	14
Folds: 1...5....10	CVF = .6787364		
Grid value 41:	lambda = .0071831	no. of nonzero coef. =	15
Folds: 1...5....10	CVF = .6801017		
... cross-validation complete ... minimum found			

Lasso logit model	No. of obs	=	358
	No. of covariates	=	31
Selection: Cross-validation	No. of CV folds	=	10

ID	Description	lambda	No. of nonzero coef.	Out-of- sample dev. ratio	CV mean deviance
1	first lambda	.2968057	0	0.0046	1.362455
36	lambda before	.0114375	13	0.5006	.6772458
* 37	selected lambda	.0104214	14	0.5007	.6771735
38	lambda after	.0094956	14	0.5005	.6774453
41	last lambda	.0071831	15	0.4985	.6801017

* lambda selected by cross-validation.

Our selected lambda is indicated by an asterisk, and can be seen visually in the cross-validation plot. Here, the Lasso has selected a lambda with 14 coefficients, some of which we are going to want to prune down.



```
. lassocoeff, display(coef, standardized)
```

	active
cost	.1421825
tortilla	.43437
temp	.4847077
meat	1.072969

Legend:

- b - base level
- e - empty cell
- o - omitted

Statistical Methods

top_burrito ~ synergy + meat + meatfillings + fillings

```
. set seed 453
. bayes, burnin(5000):logit top_burrito c.(synergy meat meatfilling fillings)
Burn-in ...
Simulation ...
```

```

-
Likelihood:
  top_burrito ~ logit(xb_top_burrito)

Prior:
  {top_burrito:synergy meat meatfilling fillings _cons} ~ normal(0,10000)
(1)
-----
-
(1) Parameters are elements of the linear form xb top burrito.

```

Bayesian logistic regression MCMC iterations =
15,000

Random-walk Metropolis-Hastings sampling	Burn-in	=
5,000		
	MCMC sample size	=
10,000		
	Number of obs	=
404		
	Acceptance rate	=
.256		
	Efficiency: min	=
.03947		
		avg =
.04696		
Log marginal-likelihood = -160.34572		max =
.05357		

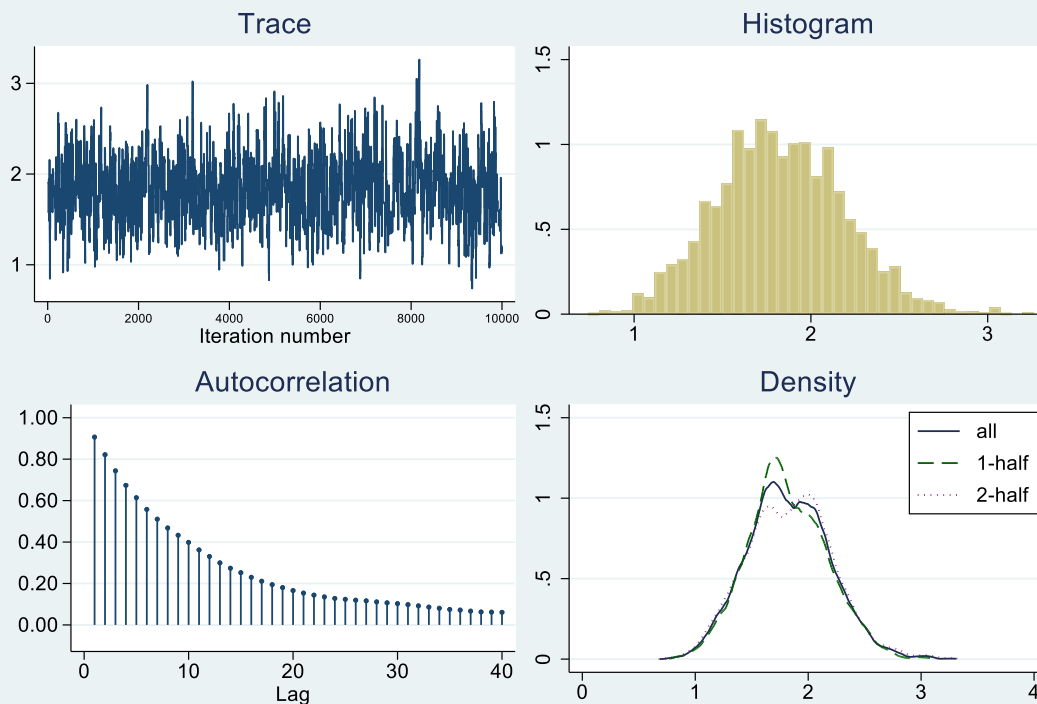
	Mean	Std. Dev.	MCSE	Median	Equal-tailed [95% Cred.
synergy	1.826937	.3628998	.017686	1.810796	1.15451
meat	1.624753	.3205945	.014894	1.612789	1.019055
meatfilling	1.051048	.2255301	.009766	1.053912	.5946789
fillings	1.203905	.3318068	.016702	1.198985	.5724701
_cons	-21.52936	2.332865	.100791	-21.5	-26.42333

Note: Default priors are used for model parameters.

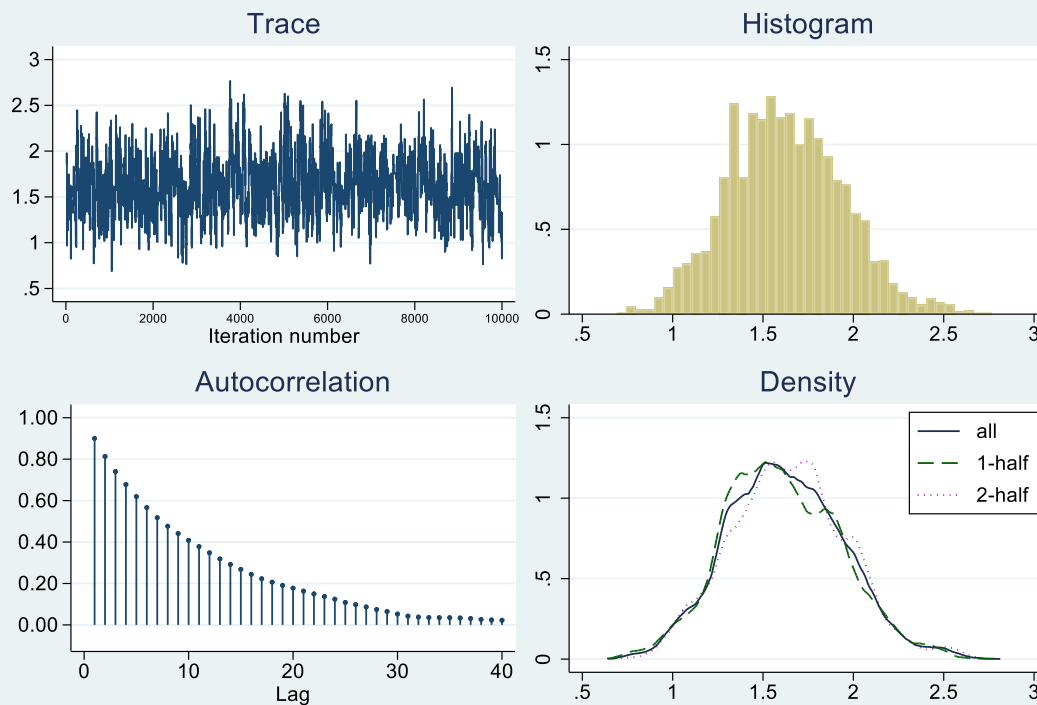
Note that we used independent normal priors in the formulation of this model—these are weakly informative at best. Additionally, I requested a longer burn-in period in the call than Stata's default in order to give the model a longer adaptation period.

Before we get into the interpretation of the model, we should perform convergence diagnostics.


top_burrito:synergy




top_burrito:meat

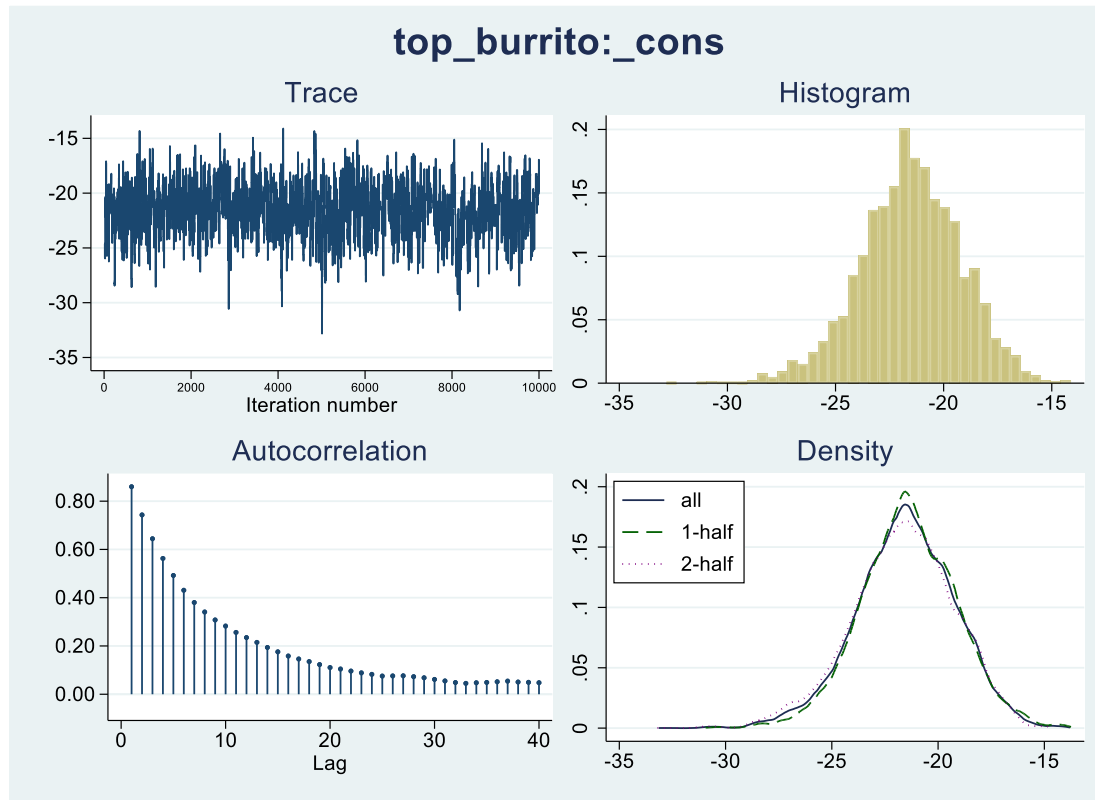


top_burrito:meatfilling



top_burrito:fillings





Overall, these convergence plots look fine. The trace is constant for each and it looks like the algorithm mixed through the posterior distribution adequately. Our autocorrelations become negligible over time, another thing we expect to see in a converged model. Additionally, the histograms are in good agreement with the normal distribution.

Referencing the model output, our acceptance rate of 0.256 is in the desired range.

Our effective samples sizes for each covariate are given below.

```
. bayesstats ess
```

```
Efficiency summaries      MCMC sample size =    10,000
                          Efficiency:  min =     .03947
                                      avg =     .04696
                                      max =     .05357
```

top_burrito	ESS	Corr. time	Efficiency
synergy	421.03	23.75	0.0421
meat	463.32	21.58	0.0463
meatfilling	533.27	18.75	0.0533
fillings	394.68	25.34	0.0395
_cons	535.72	18.67	0.0536

We would ideally like to see ESS measures closer to the MCMC size, and efficiencies closer to 10%. One solution to this problem would be to run the model with higher sample sizes, or increase the number of chains.

Note: I took this step (rerunning with 2 chains and MCMC size of 30,0000) but still had issues with efficiency. I'm not including these results here for ease of exposition.

Results

Having examined the convergence diagnostics for our model, we turn to interpretation. I'll run our model again, but request odds ratios in the command.

```
. bayes, or

Model summary
-----
-
Likelihood:
  top_burrito ~ logit(xb_top_burrito)

Prior:
  {top_burrito:synergy meat meatfilling fillings _cons} ~ normal(0,10000)
(1)
-----
-
(1) Parameters are elements of the linear form xb_top_burrito.

Bayesian logistic regression                MCMC iterations =
15,000
Random-walk Metropolis-Hastings sampling    Burn-in          =
5,000
                                           MCMC sample size =
10,000
                                           Number of obs      =
404
                                           Acceptance rate    =
.256
                                           Efficiency: min =
.03997
                                           avg =
.06036
Log marginal-likelihood = -160.34572        max =
.1211
-----
-
top_burrito |
|Odds Ratio   Std. Dev.    MCSE      Median   Equal-tailed
Interval]    [95% Cred.
-----+-----
synergy |    6.649253   2.604627   .130276   6.115314   3.17247
12.89497
```

```

      meat | 5.350071 1.808814 .083706 5.016781 2.770574
9.738832
    meatfilling | 2.934869 .678627 .029336 2.868853 1.812449
4.46952
      fillings | 3.523644 1.217598 .060496 3.31675 1.77264
6.347822
      _cons | 5.05e-09 2.81e-08 8.1e-10 4.60e-10 3.35e-12 3.43e-
08
-----

```

Note: `_cons` estimates baseline odds.

Note: Default priors are used for model parameters.

The baseline odds are given by the constant, which is close to zero. For an example of interpretation, the estimated odds of a “top burrito rating” (≥ 4) are 6.6 times as large for an increase in synergy rating of one point, where the other covariates are held constant. The 95% credible associated with this odds ratio is (3.17, 12.89).

How do we obtain predictions for our model? Because of some peculiarities of Stata, the easiest way to do so is to code the model in a slightly different (but fairly equivalent) form using a different command which requests the Random-walk Metropolis-Hastings method specifically.

```

. set seed 453

. bayesmh top_burrito c.(synergy meat meatfilling fillings),
likelihood(logit) prior({synergy} {meat} {meatfilling} {fillings} {_cons},
normal(0,10000)) burnin(5000
> ) saving(bayes, replace)

```

Burn-in ...

Simulation ...

Model summary

Likelihood:

```
top_burrito ~ logit(xb_top_burrito)
```

Prior:

```
{top_burrito:synergy meat meatfilling fillings _cons} ~ normal(0,10000)
(1)
```

(1) Parameters are elements of the linear form `xb_top_burrito`.

Bayesian logistic regression	MCMC iterations	=
15,000		
Random-walk Metropolis-Hastings sampling	Burn-in	=
5,000		
	MCMC sample size	=
10,000		
	Number of obs	=
404		

```

.256
.03947
.04696
Log marginal-likelihood = -160.34572
.05357

Acceptance rate =
Efficiency: min =
avg =
max =

```

		Mean	Std. Dev.	MCSE	Median	Equal-tailed [95% Cred. Interval]
top_burrito						
synergy		1.826937	.3628998	.017686	1.810796	1.15451
meat		1.624753	.3205945	.014894	1.612789	1.019055
meatfilling		1.051048	.2255301	.009766	1.053912	.5946789
fillings		1.203905	.3318068	.016702	1.198985	.5724701
_cons		-21.52936	2.332865	.100791	-21.5	-26.42333

file bayes.dta saved

Doing so, we can take advantage of Stata's suite of **bayesspredict** commands.

```
. bayesspredict {_ysim1}, saving(predfile, replace)
```

Computing predictions ...

```
file predfile.dta saved
file predfile.ster saved
```

```
. use predfile.dta, clear
```

```
. list _chain _index _ysim1_1 _mul_1 _frequency in 1/10
```

	_chain	_index	_ysim1_1	_mul_1	_frequency
1.	1	1	0	.22726647	1
2.	1	2	0	.22726647	1
3.	1	3	0	.22726647	1
4.	1	4	0	.22726647	1
5.	1	5	0	.22726647	1
6.	1	6	1	.22726647	1

```

7. |      1      7      0      .22726647      1 |
8. |      1      8      0      .22726647      1 |
9. |      1      9      0      .22726647      1 |
10. |     1     10      1      .22726647      1 |
+-----+

```

Focusing only on the first burrito, we see a list of the first 10 simulated predictions. This dataset of predictions is actually quite large; there are a total of 10,000 similar predicted values for each of the 423 burritos in the original data set. Essentially, the **bayesspredict** command uses samples the posterior predictive distribution for our outcome of *top_burrito* 10,000 times for each burrito.

Note that we can apply Stata's normal commands to this simulated data set. For example, let's summarize all 10,000 results for the first two burritos.

```
. bayesstats summary {_ysim[1] _ysim[2] _mu[1] _mu[2]} using predfile
```

```
Posterior summary statistics                                MCMC sample size =
10,000
```

```

-----
-
      |
      |      Mean   Std. Dev.   MCSE   Median   Equal-tailed
Interval]      [95% Cred.
-----+-----
-
      |
1  _ysim1_1 |      .2881   .4529005   .00472      0      0
0  _ysim1_2 |      .0007   .0264496   .000264     0      0
      |
      |      .2874842   .0643705   .002389   .2863605   .1678432
      |      .4182073
      |      .0005911   .0005618   .000022   .0004209   .000066
      |      .0020245
-----
-

```

There is much more that Stata can do for postestimation after Bayesian analysis. For example, we could perform out of sample predictions. However, in the interest of time I'll leave this here.

[Comparison to Previous Analyses](#)

Our main point of comparison is Scott Cole's principal components analysis, referenced previously in the "Background" section. This analysis found that the ratings for fillings, meat, meat/filling ratio, and uniformity were the most important (note: he excluded synergy rating as a predictor due to potential difficulty disentangling it from the outcome).

Our Bayesian model is largely in agreement with this analysis, with pertinent differences being our inclusion of synergy and lack of inclusion of uniformity, as it did not seem to be important in our variable selection process.

Limitations/Statistical Issues

The main limitation of this analysis is use of noninformative prior probabilities. Bayesian models using noninformative priors like ours may be less subjective, but this comes at the price of loss of customizability and practical application of probability knowledge. Arguably, we lose the chief advantages of the Bayesian approach in the first place.

```
. set seed 453

. bayes, burnin(5000): logit top_burrito c.(synergy meat meatfilling
fillings)

Burn-in ...
Simulation ...

Model summary
-----
-
Likelihood:
    top_burrito ~ logit(xb_top_burrito)

Prior:
    {top_burrito:synergy meat meatfilling fillings _cons} ~ normal(0,10000)
(1)
-----
-
(1) Parameters are elements of the linear form xb_top_burrito.

Bayesian logistic regression                                MCMC iterations =
15,000
Random-walk Metropolis-Hastings sampling                    Burn-in          =
5,000
                                                                MCMC sample size =
10,000
                                                                Number of obs     =
404
                                                                Acceptance rate   =
.256
                                                                Efficiency: min =
.03947
                                                                avg =
.04696
Log marginal-likelihood = -160.34572                          max =
.05357
-----
-

```

	Mean	Std. Dev.	MCSE	Median	Equal-tailed [95% Cred. Interval]
top_burrito					
synergy	1.826937	.3628998	.017686	1.810796	1.15451

```
-----
-
2.556837
```



```

      meat | 1.624753 .3205945 .014894 1.612789 1.019055
2.276121
    meatfilling | 1.051048 .2255301 .009766 1.053912 .5946789
1.497281
      fillings | 1.203905 .3318068 .016702 1.198985 .5724701
1.848112
      _cons | -21.52936 2.332865 .100791 -21.5 -26.42333 -
17.18702
-----

```

Note: Default priors are used for model parameters.

```
. logit top_burrito c.(synergy meat meatfilling fillings)
```

```

Iteration 0: log likelihood = -275.85367
Iteration 1: log likelihood = -141.97749
Iteration 2: log likelihood = -130.84006
Iteration 3: log likelihood = -130.56487
Iteration 4: log likelihood = -130.56421
Iteration 5: log likelihood = -130.56421

```

```

Logistic regression                                Number of obs      =
404                                                  LR chi2(4)         =
290.58                                              Prob > chi2        =
0.0000                                              Pseudo R2          =
Log likelihood = -130.56421
0.5267

```

```

-----
-
top_burrito |      Coef.   Std. Err.      z    P>|z|     [95% Conf.
Interval]
-----+-----
-
      synergy | 1.765252   .356725    4.95   0.000    1.066084
2.46442
      meat | 1.566101   .3074037    5.09   0.000    .9636007
2.168601
    meatfilling | 1.024583   .2191053    4.68   0.000    .5951446
1.454022
      fillings | 1.168341   .3259696    3.58   0.000    .5294522
1.80723
      _cons | -20.84379  2.233457   -9.33   0.000   -25.22128   -
16.46629
-----
-

```

Looking at our output, now back on the log odds scale, we note that our estimated coefficients and confidence/credible intervals are quite similar between the two approaches. This would not be the case if we had different priors, perhaps from a pilot study or subject matter expertise. This would have likely helped us with the efficiency problems as well.

Another limitation is the logistic design of the outcome variable. It would have been better to keep the outcome as a multilevel, ordinal rating (I had convergence problems when I tried this). An additional limitation, in my opinion, is the lack of inclusion of any indicator variables corresponding to ingredients. It would have been nice to discover, for example, that the presence of guacamole or bacon or something else helped elevate a burrito's rating independent of other factors.

Future Directions

Research, as properly conceived, is dynamic— it is not a static event and cannot be contained in a single study or analysis. Having seen some of the things that factor into a quality burrito, and in the process learning a whole lot about Bayesian methods, I am not content to stop here. I suspect that it will take years of further study and the ingestion of hundreds of burritos for me to really get to the state of zen-like knowledge and culinary satisfaction I am hoping to achieve. However, I'm a scientist, and I'm willing to take the time, and the pounds, to see this labor of burrito love to completion.